

house-price

May 8, 2023

```
[13]: from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras import Sequential
from sklearn.metrics import r2_score
from matplotlib import pyplot as plt
import pandas as pd
import numpy as np
```

```
[14]: df = pd.read_csv(r'C:\Users\nithy\Downloads\Compressed\House Price India.csv')
df.drop(['id', 'Date'], axis=1, inplace=True)
df.head()
```

```
[14]:
```

	number of bedrooms	number of bathrooms	living area	lot area	\
0	5	2.50	3650	9050	
1	4	2.50	2920	4000	
2	5	2.75	2910	9480	
3	4	2.50	3310	42998	
4	3	2.00	2710	4500	

	number of floors	waterfront present	number of views	\
0	2.0	0	4	
1	1.5	0	0	
2	1.5	0	0	
3	2.0	0	0	
4	1.5	0	0	

	condition of the house	grade of the house	\
0	5	10	
1	5	8	
2	3	8	
3	3	9	
4	4	8	

	Area of the house(excluding basement)	...	Built Year	Renovation Year	\
0	3370	...	1921	0	
1	1910	...	1909	0	

2		2910	...	1939	0
3		3310	...	2001	0
4		1880	...	1929	0

	Postal Code	Lattitude	Longitude	living_area_renov	lot_area_renov	\
0	122003	52.8645	-114.557	2880	5400	
1	122004	52.8878	-114.470	2470	4000	
2	122004	52.8852	-114.468	2940	6600	
3	122005	52.9532	-114.321	3350	42847	
4	122006	52.9047	-114.485	2060	4500	

	Number of schools nearby	Distance from the airport	Price
0	2	58	2380000
1	2	51	1400000
2	1	53	1200000
3	3	76	838000
4	1	51	805000

[5 rows x 21 columns]

[15]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   number of bedrooms                   14620 non-null  int64
1   number of bathrooms                  14620 non-null  float64
2   living area                          14620 non-null  int64
3   lot area                             14620 non-null  int64
4   number of floors                     14620 non-null  float64
5   waterfront present                  14620 non-null  int64
6   number of views                      14620 non-null  int64
7   condition of the house               14620 non-null  int64
8   grade of the house                  14620 non-null  int64
9   Area of the house(excluding basement) 14620 non-null  int64
10  Area of the basement                 14620 non-null  int64
11  Built Year                           14620 non-null  int64
12  Renovation Year                      14620 non-null  int64
13  Postal Code                          14620 non-null  int64
14  Lattitude                           14620 non-null  float64
15  Longitude                            14620 non-null  float64
16  living_area_renov                    14620 non-null  int64
17  lot_area_renov                       14620 non-null  int64
18  Number of schools nearby              14620 non-null  int64
19  Distance from the airport            14620 non-null  int64
```

```
20 Price 14620 non-null int64
dtypes: float64(4), int64(17)
memory usage: 2.3 MB
```

```
[16]: df.isna().sum()
```

```
[16]: number of bedrooms      0
      number of bathrooms    0
      living area            0
      lot area               0
      number of floors       0
      waterfront present     0
      number of views        0
      condition of the house 0
      grade of the house     0
      Area of the house(excluding basement) 0
      Area of the basement   0
      Built Year             0
      Renovation Year        0
      Postal Code            0
      Lattitude              0
      Longitude              0
      living_area_renov      0
      lot_area_renov         0
      Number of schools nearby 0
      Distance from the airport 0
      Price                  0
      dtype: int64
```

```
[17]: X = df.drop('Price', axis=1)
      y = df['Price']
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=12)
```

```
[19]: print(f'\n shape of X_train - {X_train.shape}\n')
      print(f' shape of X_test - {X_test.shape}\n')
      print(f' shape of y_train - {y_train.shape}\n')
      print(f' shape of y_test - {y_test.shape}\n')
```

```
shape of X_train - (11696, 20)
```

```
shape of X_test - (2924, 20)
```

```
shape of y_train - (11696,)
```

```
shape of y_test - (2924,)
```

```
[30]: number_of_features = len(X.columns)

model = Sequential()

model.add(layer=Input(shape=number_of_features))

model.add(layer=Dense(units=32, activation='relu'))

model.add(layer=Dense(units=64, activation='relu'))

model.add(layer=Dense(units=128, activation='relu'))

model.add(layer=Dense(units=256, activation='relu'))

model.add(layer=Dense(units=512, activation='relu'))

model.add(layer=Dense(units=1024, activation='relu'))

model.add(layer=Dense(units=512, activation='relu'))

model.add(layer=Dense(units=256, activation='relu'))

model.add(layer=Dense(units=128, activation='relu'))

model.add(layer=Dense(units=64, activation='relu'))

model.add(layer=Dense(units=32, activation='relu'))

model.add(layer=Dense(units=16, activation='relu'))

model.add(layer=Dense(units=1, activation='linear'))
```

```
[31]: adam = Adam(learning_rate=7e-5)

model.compile(optimizer=adam, loss='mse', metrics=['mae', 'mape'])
```

```
[33]: history = model.fit(X_train, y_train, epochs=10)
```

```
Epoch 1/10
366/366 [=====] - 7s 20ms/step - loss: 21744080896.0000
- mae: 106028.5234 - mape: 23.9161
Epoch 2/10
366/366 [=====] - 8s 22ms/step - loss: 21374562304.0000
- mae: 105025.5625 - mape: 23.7174
Epoch 3/10
366/366 [=====] - 7s 20ms/step - loss: 19224690688.0000
```

```

- mae: 100039.5000 - mape: 22.9970
Epoch 4/10
366/366 [=====] - 8s 23ms/step - loss: 18683500544.0000
- mae: 99278.1172 - mape: 22.8519
Epoch 5/10
366/366 [=====] - 8s 22ms/step - loss: 18234394624.0000
- mae: 98242.8594 - mape: 22.7213
Epoch 6/10
366/366 [=====] - 8s 22ms/step - loss: 20518877184.0000
- mae: 102074.7656 - mape: 23.2780
Epoch 7/10
366/366 [=====] - 7s 19ms/step - loss: 22670522368.0000
- mae: 106679.6953 - mape: 23.9765
Epoch 8/10
366/366 [=====] - 7s 18ms/step - loss: 18700814336.0000
- mae: 99613.9766 - mape: 22.8452
Epoch 9/10
366/366 [=====] - 6s 17ms/step - loss: 18543321088.0000
- mae: 98677.9766 - mape: 22.7990
Epoch 10/10
366/366 [=====] - 7s 18ms/step - loss: 19590918144.0000
- mae: 101507.4219 - mape: 23.2146

```

```
[34]: y_pred = model.predict(X_test)
```

```
92/92 [=====] - 1s 6ms/step
```

```
[35]: y_pred
```

```
[35]: array([[387830.6 ],
           [445728.56],
           [370873.7 ],
           ...,
           [281568.03],
           [358895.47],
           [703242.4 ]], dtype=float32)
```

```
[36]: y_test
```

```
[36]: 12149    640000
      13581    650000
      11595    325000
      2769    373000
      7393    355000
      ...
      7362    497000
      11132    400000
      142     366750
```

```
1405      276000
6184      569000
Name: Price, Length: 2924, dtype: int64
```

```
[37]: pd.DataFrame({'Actual Value':y_test.values.flatten(), 'Predicted Value':y_pred.
↪flatten()})
```

```
[37]:
```

	Actual Value	Predicted Value
0	640000	387830.59375
1	650000	445728.56250
2	325000	370873.68750
3	373000	476821.93750
4	355000	465932.37500
...
2919	497000	604781.50000
2920	400000	339529.09375
2921	366750	281568.03125
2922	276000	358895.46875
2923	569000	703242.37500

```
[2924 rows x 2 columns]
```

```
[38]: r2_score(y_pred=y_pred, y_true=y_test)
```

```
[38]: 0.5681449923201896
```

```
[ ]:
```